



## Administrative Notes- Jan 26, 2023

- Jan 27: Assignment 1 due
- Assignment 2 available on Canvas
- Assignment 3 released on Friday night
  - You are allowed to complete this in pairs (more instructions on how to sign up as a pair in the assignment description released on Friday night)
- You will need a physical student card for the class so [please get one](#)



# Review

## Student

sid	name	address	phone	major
99111120	K. Jones	1234 W. 12 <sup>th</sup> Ave., Van	889-4444	CPSC
92001200	S. Selvarajah	2020 E. 18 <sup>th</sup> St., Van	409-2222	MATH
94001020	A. Alberty	2020 E. 18 <sup>th</sup> St., Van	222-2222	FREN
94001150	J. Wang	null	null	null

Student(sid: integer, name: string, address: string, phone: string, major: string)

Or without the domains:

Student (sid, name, address, phone, major)



# Review

- The statement on the right creates the Student relation
  - the type (**domain**) of each attribute is specified and enforced when tuples are added or modified
- The statement on right creates Grade information about courses that a student takes

```
CREATE TABLE Student
(sid      INTEGER,
 name    CHAR(20),
 address CHAR(30),
 phone   CHAR(13),
 major   CHAR(4))
```

```
CREATE TABLE Grade
(sid      INTEGER,
 dept     CHAR(4),
 course#  CHAR(3),
 mark     INTEGER)
```



# Keys Constraints (for Relations)

- Similar to those for entity sets in the ER model
- One or more attributes in a relation form a key (or candidate key) for a relation, where S is the set of all attributes in the key, if:
  1. No distinct tuples can have the same values for all attributes in the key, and
  2. No subset of S is itself a key (according to (1)).(If such a subset exists, then S is a superkey and not a key.)
- One of the possible keys is chosen (by the DBA) to be the primary key (PK).

CREATE TABLE Student

```
(sid      INTEGER PRIMARY KEY,  
name     CHAR(20),  
address  CHAR(30),  
phone    CHAR(13),  
major    CHAR(4))
```



# Keys Constraints in SQL

- A **PRIMARY KEY** constraint specifies a table's primary key
  - values for primary key must be unique
  - a primary key attribute cannot be *null*
- Other keys are specified using the **UNIQUE** constraint
  - values for a group of attributes must be unique (if they are not null)
  - these attributes can be *null*
- Key constraints are checked when
  - new values are inserted
  - values are modified

# Clicker Question

Does the constraint PRIMARY KEY(Dept, Course#) hold for this instance?

Course#	Dept	Title
100	CPSC	Computational Thinking
100	MATH	Differential Calculus with Applications to Physical Sciences and Engineering

A. Yes

B. No

C. It depends

# Clicker Question

Does the constraint PRIMARY KEY(Dept, Course#) hold for this instance?

Dept	Course #	Term	Section
CPSC	368	2021W2	201
CPSC	368	2021W2	202

A. Yes

B. No

C. It depends



# Keys Constraints in SQL (cont')

(Ex.1- Normal) “For a given student and course, there is a single grade.”

VS.

(Ex.2 - Silly) “Students can take a course once, and receive a single grade for that course; further, no two students in a course receive the same grade.”

```
CREATE TABLE Grade
(sid      INTEGER,
dept     CHAR(4),
course#  CHAR(3),
mark     INTEGER,
PRIMARY KEY (sid,dept,course#) )
```

```
CREATE TABLE Grade2
(sid      INTEGER,
dept     CHAR(4),
course#  CHAR(3),
mark     CHAR(2),
PRIMARY KEY (sid,dept,course#),
UNIQUE (dept,course#,mark) )
```





# Keys Constraints in SQL (cont')

For single attribute keys, can also be declared on the same line as the attribute.

```
CREATE TABLE Student  
  (sid      INTEGER PRIMARY KEY,  
   name     CHAR(20),  
   address  CHAR(30),  
   phone    CHAR(13),  
   major    CHAR(4))
```



# Foreign Keys Constraints

- **Foreign key** : Set of attributes in one relation used to 'reference' a tuple in another relation.
  - **Must correspond to the primary key of the other relation.**
  - Like a 'logical pointer'.
- E.g.: Grade(*sid*, *dept*, *course#*, *grade*)
  - *sid* is a foreign key referring to **Student**:
  - (*dept*, *course#*) is a foreign key referring to **Course**
- **Referential integrity**: All foreign keys reference existing entities.
  - i.e. there are no dangling references
  - all foreign key constraints are enforced



# Foreign Keys in SQL

Only students listed in the Student relation should be allowed to have grades for courses that are listed in the Course relation.

CREATE TABLE Grade

(sid INTEGER, dept CHAR(4), course# CHAR(3), mark INTEGER,  
**PRIMARY KEY** (sid,dept,course#),  
**FOREIGN KEY** (sid) **REFERENCES** Student,  
**FOREIGN KEY** (dept, course#) **REFERENCES** Course(dept, cnum))


Sometimes you can not specify which attributes are referenced, but in this case they are needed. Never hurts to include them!

Grade

sid	dept	course#	mark
53666	CPSC	101	80
53666	RELG	100	45
53650	MATH	200	null
53666	HIST	201	60

Student

sid	name	address	Phone	major
53666	G. Jones	....	...	...
53688	J. Smith	....	...	...
53650	G. Smith	....	...	...





# Enforcing Referential Integrity

```
CREATE TABLE Grade (  
    sid INTEGER,  
    dept CHAR(4),  
    course# CHAR(3),  
    mark INTEGER,  
    PRIMARY KEY (sid,dept,course#),  
    FOREIGN KEY (sid) REFERENCES  
        Student,  
    FOREIGN KEY (dept, course#)  
        REFERENCES Course(dept, cnum)  
    )
```

## Grade

sid	dept	cnum	grade
2	CPSC	304	90
2	MATH	221	90
2	EPSE	223	90
1	MUSC	103	90


# Enforcing Referential Integrity

Student

sid	name	...
1	Blossom	...
2	Buttercup	...
3	Bubbles	...
4	Blossom	

Grade

sid	dept	cnum	grade
2	CPSC	304	90
2	MATH	221	90
2	EPSE	223	90
1	MUSC	103	90



A foreign key is a set of attributes in one relation (e.g., Grades.sid) used to 'reference' a tuple in another relation (e.g., Students.sid).



# Enforcing Referential Integrity

- *sid* in Grade is a foreign key that references Student.
- What should be done if a Grade tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a **Student tuple** is deleted?
  - Also delete all Grade tuples that refer to it?
  - Disallow deletion of this particular Student tuple?
  - Set *sid* in Grade tuples that refer to it, to *null*, (the special value denoting '*unknown*' or '*inapplicable*'.)
    - Problem if *sid* is part of the primary key
  - Set *sid* in Grade tuples that refer to it, to a *default sid*.
- Similar if primary key of a Student tuple is updated



# Referential Integrity in SQL/92

- SQL/92 supports all 4 options on deletes and updates.
    - Default is **NO ACTION** (*delete/update is rejected*)
    - **CASCADE** (also updates/deletes all tuples that refer to the updated/deleted tuple)
    - **SET NULL / SET DEFAULT** (referencing tuple value is set to the default foreign key value )
- ```
CREATE TABLE Grade
(sid CHAR(8), dept CHAR(4),
course# CHAR(3), mark INTEGER,
PRIMARY KEY (sid,dept,course#),
FOREIGN KEY (sid)
REFERENCES Student(sid)
ON DELETE CASCADE
ON UPDATE CASCADE
FOREIGN KEY (dept, course#)
REFERENCES
Course(dept,course#)
ON DELETE SET DEFAULT
ON UPDATE CASCADE );
```



# Clicker Question

Consider the following table definition.

```
CREATE TABLE BMW ( bid INTEGER, sid INTEGER, ...  
                    PRIMARY KEY (bid),  
                    FOREIGN KEY (sid) REFERENCES STUDENTS  
                    ON DELETE CASCADE);
```

If  $\text{bid} = 1000$  and  $\text{sid} = 5678$  for a row in Table BMW, choose the best answer.

- A. If the row for  $\text{sid}$  value 5678 in STUDENTS is deleted, then the row with  $\text{bid} = 1000$  in BMW is automatically deleted.
- B. If a row with  $\text{sid}$  value 5678 in BMW is deleted, then the row with  $\text{sid}=5678$  in STUDENTS is automatically deleted.
- C. Both of the above.





# Clicker Question

Consider the following table definition.

```
CREATE TABLE BMW ( bid INTEGER, sid INTEGER, ...  
    PRIMARY KEY (bid),  
    FOREIGN KEY (sid) REFERENCES STUDENTS  
    ON DELETE CASCADE);
```

- A. If the row for sid value 5678 in STUDENTS is deleted, then the row with bid = 1000 in BMW is automatically deleted.
- B. If a row with sid value 5678 in BMW is deleted, then the row with sid=5678 in STUDENTS is automatically deleted.
- C. Both of the above.

BMW

| bid  | Sid  |
|------|------|
| 1000 | 5678 |

Student

| sid  | name  | Address |
|------|-------|---------|
| 5678 | James | Null    |






# Where do ICs Come From?

- ICs are based upon the real-world semantics being described (in the database relations).
- We *can* check a database instance to verify an IC, but we *cannot* tell the ICs by looking at the instance.
  - For example, even if all student names differ, we cannot assume that name is a key.
  - An IC is a statement about *all possible* instances.
- All constraints must be identified during the conceptual design.
- Some constraints can be explicitly specified in the conceptual model
  - Key and foreign key ICs are shown on ER diagrams.
- Others are written in a more general language.




# Logical DB Design: ER to Relational

- Each entity set is mapped to a table.
  - Entity attributes become table attributes
  - Entity keys become table keys



```
CREATE TABLE MoviePeople  
(ID CHAR(11),  
Name CHAR(20),  
BirthDate DATE,  
PRIMARY KEY (ID))
```

# Relationship Sets to Tables




- A relationship set id is mapped to a single relation (table).
- Simple case: relationship has no constraints (i.e. many-to-many)
- In this case, attributes of the table must include:
  - Keys for each participating entity set as foreign keys.
    - This is a **key** for the relation.
  - All descriptive attributes.

```

CREATE TABLE WorkOn(
    MPID    CHAR(11),
    MID     INTEGER,
    Role    CHAR(20),
    PRIMARY KEY (MPID, MID),
    FOREIGN KEY (MPID)
        REFERENCES MoviePeople,
    FOREIGN KEY (MID)
        REFERENCES Movies)
    
```

# Example: Many to Many Relationships




Scrooge McDuck  
has worked on  
two movies



Pua has worked  
on one movie  
(Moana)



# Example: Many to Many Relationships



Can we reduce redundancy by combining these tables in any way?



MoviePeople

| <u>ID</u> | Name    | Birthdate |
|-----------|---------|-----------|
| 1         | Scrooge | ...       |
| 2         | Pua     | ...       |

WorkOn

| <u>MPID</u> | <u>MID</u> | Role |
|-------------|------------|------|
| 1           | 1          | ...  |
| 1           | 2          | ...  |
| 2           | 3          | ...  |

Movies

| <u>ID</u> | Title             |
|-----------|-------------------|
| 1         | Ducktales         |
| 2         | A Christmas Carol |
| 3         | Moana             |

# Example: Many to Many Relationships

## MoviePeople

| <u>ID</u> | Name    | Birthdate |
|-----------|---------|-----------|
| 1         | Scrooge | ...       |
| 2         | Pua     | ...       |

## Movies

| <u>ID</u> | Title             |
|-----------|-------------------|
| 1         | Ducktales         |
| 2         | A Christmas Carol |
| 3         | Moana             |

## WorkOn

| <u>MPID</u> | <u>MID</u> | Role |
|-------------|------------|------|
| 1           | 1          | ...  |
| 1           | 2          | ...  |
| 2           | 3          | ...  |

Can we integrate the information in WorkOn into MoviePeople?



# Example: Many to Many Relationships

MoviePeople

| <u>ID</u> | Name    | Birthdate | MID |
|-----------|---------|-----------|-----|
| 1         | Scrooge | ...       |     |
| 2         | Pua     | ...       |     |

Movies

| <u>ID</u> | Title             |
|-----------|-------------------|
| 1         | Ducktales         |
| 2         | A Christmas Carol |
| 3         | Moana             |

WorkOn

| <u>MPID</u> | <u>MID</u> | Role |
|-------------|------------|------|
| 1           | 1          | ...  |
| 1           | 2          | ...  |
| 2           | 3          | ...  |

Do we put MID 1 or MID 2? We can't have both in the column (i.e., we can't store a list there)

Not much we can do in terms of reducing tables



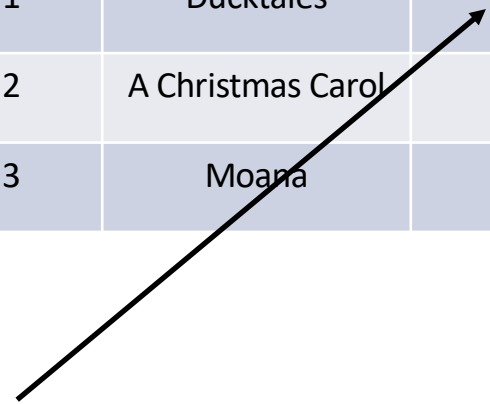
# Example: Many to Many Relationships

MoviePeople

| <u>ID</u> | Name    | Birthdate |
|-----------|---------|-----------|
| 1         | Scrooge | ...       |
| 2         | Pua     | ...       |

Movies

| <u>ID</u> | Title             | <u>MPID</u> |
|-----------|-------------------|-------------|
| 1         | Ducktales         |             |
| 2         | A Christmas Carol |             |
| 3         | Moana             |             |



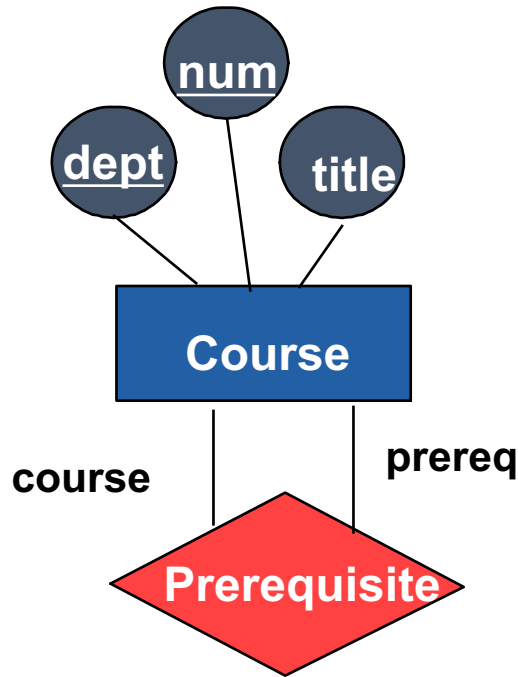
WorkOn

| <u>MPID</u> | <u>MID</u> | Role |
|-------------|------------|------|
| 1           | 1          | ...  |
| 1           | 2          | ...  |
| 2           | 3          | ...  |

What if we have more than one person work on this movie?  
Same problem as before!

Not much we can do in terms of reducing tables

# Relationship Sets to Tables (cont')



- In some cases, we need to use the roles:

```

CREATE TABLE Prerequisite(
  course_dept  CHAR(4),
  course_num   CHAR(3),
  prereq_dept  CHAR(4),
  prereq_num   CHAR(3),
  PRIMARY KEY (course_dept, course_num,
               prereq_dept, prereq_num),
  FOREIGN KEY (course_dept, course_num)
    REFERENCES Course(dept, num),
  FOREIGN KEY (prereq_dept, prereq_num)
    REFERENCES Course(dept, num))
    
```



# To motivate examples on upcoming slides, let's talk about getting a PhD

- PhD students all have to have advisors, all of whom also have had advisors (etc.)
- There exist databases where you can go back hundreds of years to see people's academic lineage
- Out of curiosity, and to make this more interesting for you (you're welcome), you can use <https://www.genealogy.math.ndsu.nodak.edu/> to look this information up



# Rachel's Academic Genealogy



- Rachel Pottinger
- Phil Bernstein
- Catriel Beeri
- Eli Shamir
- Shmuel Agmon
- Szolem Mandelbroit
- Jacques Salomon Hadamard
- C Emile (Charles) Picard
- Gaston Darboux
- Michel Chasles
- Simeon Denis Poisson
- Joseph Louis Lagrange
- Leonhard Euler
- Johann Bernoulli
- Jacob Bernoulli
- Peter Werenfels
- Theodor Zwinger, Jr.
- Sebastian Beck
- Johann Jacob Grynaeus
- Simon Sulzer
- Wolfgang Fabricius Capito
- Desiderious Erasmus
- Jan Standonck



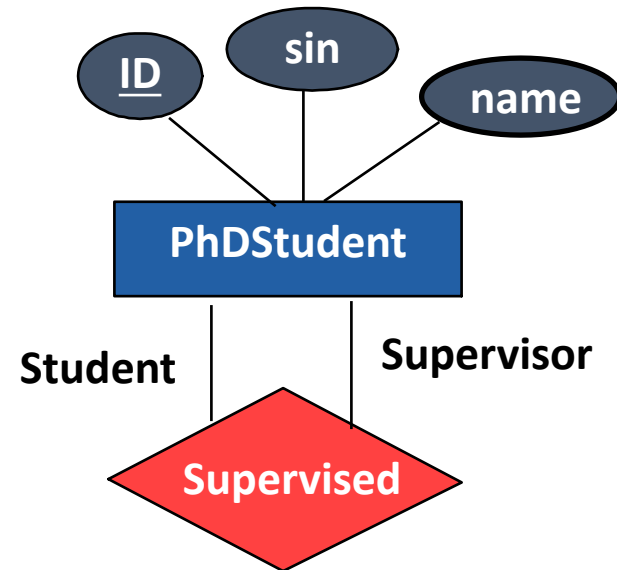
# One possible partial representation of this data is

```
CREATE TABLE PhDStudent(  
    id INT,  
    sin INT,  
    name CHAR(20),  
    advisorID INT);
```

| id | sin  | name                | AdvisorID |
|----|------|---------------------|-----------|
| 1  | Null | Jan Standonck       | Null      |
| 2  | Null | Desiderious Erasmus | 1         |

# One possible partial representation of this data is

```
CREATE TABLE PhDStudent(
    id INT,
    sin INT,
    name CHAR(20),
    advisorID INT);
```



| id | sin  | name                | AdvisorID |
|----|------|---------------------|-----------|
| 1  | Null | Jan Standonck       | Null      |
| 2  | Null | Desiderious Erasmus | 1         |



# Self Referencing Relations

**Goal:** have Advisor be foreign key reference for **same table** PhDstudent

| id | sin  | name                | AdvisorID |
|----|------|---------------------|-----------|
| 1  | Null | Jan Standonck       | Null      |
| 2  | Null | Desiderious Erasmus | 1         |

Could a foreign key be null?

For **referential integrity** to hold in a relational database, any field in a table that is declared a foreign key should contain either a **null value**, or only values from a parent table's **primary key**.



# Clicker Question





# Clicker Question

Consider the table definition: CREATE TABLE PhDStudent(  
(1) id INT,  
Numbers denote lines only -> (2) sin INT,  
(3) name CHAR(20),  
(4) advisorID INT);

Which of the following is **not** a legal addition?

- A. Add UNIQUE just before the commas on lines (2) and (3) and add PRIMARY KEY just before the comma on line (1).
- B. Add PRIMARY KEY just before the commas on lines (1) and (2).
- C. Add UNIQUE just before the comma on line (1), and add PRIMARY KEY just before the comma on line (2).
- D. All are legal
- E. None are legal



# Reasoning

- B is not legal because it attempts to create two primary keys:

```
CREATE TABLE PhDStudent(  
  id INT PRIMARY KEY,  
  sin INT PRIMARY KEY,  
  name CHAR(20),  
  advisorID INT);
```

- Creating a complex primary key that consisting of the combination of id and sin, would be done as follows:

```
CREATE TABLE PhDStudent(  
  id INT,  
  sin INT,  
  name CHAR(20),  
  advisorID INT,  
  PRIMARY KEY (id, sin));
```

- Note, that this is a **terrible** idea because each of sin and id are keys by themselves



# Clicker Question

Consider the table definition:


```
CREATE TABLE PhDStudent(  (1) id INT,  
Numbers denote lines only -> (2) sin INT,  
                             (3) name CHAR(20),  
                             (4) advisorID INT);
```

**Goal:** have advisorID be foreign key reference for same table PhDStudent.

Which of the following is not legal? (does not have to achieve all goals)


- A. Add FOREIGN KEY (advisorID) REFERENCES PhDStudent(id) before the ) on line (4).
- B. Add PRIMARY KEY just before the comma on lines (1) and (2), and add REFERENCES PhDStudent(id) before the ) on line (4).
- C. Add PRIMARY KEY just before the comma on line (1), add UNIQUE just before the comma on line (2), and add FOREIGN KEY REFERENCES PhDStudent(sin) before the ) on line (4).
- D. All are legal
- E. None are legal

# Relationship Sets with Key Constraints



- Each movie has at most one director, according to the key constraint on Direct.
- How can we take advantage of this?


# For the one to many case, combine the many side with the relationship



I know who the MoviePerson is if I know the movie!

I can remove some redundancy in my design.

# For the one to many case, combine the many side with the relationship




MoviePeople

| ID | Name | Birth Date |
|----|------|------------|
| 1  |      |            |
| 2  |      |            |

Movies

| ID | MoviePerson ID | Title |
|----|----------------|-------|
| 1  | 1              |       |
| 2  | 1              |       |
| 3  | 2              |       |


# For the one to many case, combine the many side with the relationship



James  
Cameron



You



The Life  
of a CS  
Student

# For the one to many case, combine the many side with the relationship

MoviePeople

| <u>ID</u> | Name          | Birthdate |
|-----------|---------------|-----------|
| 1         | James Cameron | ...       |
| 2         | You           | ...       |

Movies

| <u>ID</u> | Title                    |
|-----------|--------------------------|
| 1         | Avatar                   |
| 2         | Titanic                  |
| 3         | The Life of a CS Student |

Direct

| <u>MPID</u> | <u>MID</u> |
|-------------|------------|
| 1           | 1          |
| 1           | 2          |
| 2           | 3          |

Can we reduce redundancy by integrating Direct into MoviePeople?





# For the one to many case, combine the many side with the relationship

MoviePeople

| <u>ID</u> | Name          | Birthdate | Directed<br>MID |
|-----------|---------------|-----------|-----------------|
| 1         | James Cameron | ...       |                 |
| 2         | You           | ...       |                 |

Movies

| <u>ID</u> | Title                    |
|-----------|--------------------------|
| 1         | Avatar                   |
| 2         | Titanic                  |
| 3         | The Life of a CS Student |

Direct

| <u>MPID</u> | <u>MID</u> |
|-------------|------------|
| 1           | 1          |
| 1           | 2          |
| 2           | 3          |

Same issue as before.  
We can't have  
multiple values here.



# For the one to many case, combine the many side with the relationship

MoviePeople

| <u>ID</u> | Name          | Birthdate |
|-----------|---------------|-----------|
| 1         | James Cameron | ...       |
| 2         | You           | ...       |

Movies

| <u>ID</u> | Title                    |
|-----------|--------------------------|
| 1         | Avatar                   |
| 2         | Titanic                  |
| 3         | The Life of a CS Student |

Direct

| <u>MPID</u> | <u>MID</u> |
|-------------|------------|
| 1           | 1          |
| 1           | 2          |
| 2           | 3          |

Can we integrate  
Direct into Movies?



# For the one to many case, combine the many side with the relationship

MoviePeople

| <u>ID</u> | Name          | Birthdate |
|-----------|---------------|-----------|
| 1         | James Cameron | ...       |
| 2         | You           | ...       |

Movies

| <u>ID</u> | Title                    | Director-MPID |
|-----------|--------------------------|---------------|
| 1         | Avatar                   | 1             |
| 2         | Titanic                  | 1             |
| 3         | The Life of a CS Student | 2             |



Will there ever be two different directors for the same movie?


No! (because of our one to many relationship)



# We'd covered basic translating of ER to relational

- Short version: everything's a table
- Slightly longer version: in many to many relationships, create one table per entity and one table per relationship. Link the two by foreign keys

# For the one to many case, combine the many side with the relationship



MoviePeople

| <u>ID</u> | Name          | Birthdate |
|-----------|---------------|-----------|
| 1         | James Cameron | ...       |
| 2         | You           | ...       |

Movies

| <u>ID</u> | Title                    | Director-MPID |
|-----------|--------------------------|---------------|
| 1         | Avatar                   | 1             |
| 2         | Titanic                  | 1             |
| 3         | The Life of a CS Student | 2             |



# Translating ER Diagrams with Key Constraints

- Method 1 (unsatisfactory):
  - Create a separate table for Direct:
  - Note that **MID** is the key now!
  - Create separate tables for MoviePeople and Movies.


```
CREATE TABLE Direct(  
  MPID CHAR(11),  
  MID INTEGER,  
  PRIMARY KEY (MID),  
  FOREIGN KEY (MPID) REFERENCES  
MoviePeople,  
  FOREIGN KEY (MID) REFERENCES Movies)
```

- Method 2 (better)
  - Since each movie has a unique director, we can **combine Direct and Movies into one table.**
  - Create another table for MoviePeople
  - Must have on delete and on update in this case!

```
CREATE TABLE Directed_Movie(  
  MID INTEGER,  
  title CHAR(20),  
  MPID CHAR(11),  
  PRIMARY KEY (MID),  
  FOREIGN KEY (MPID) REFERENCES  
MoviePeople  
  ON DELETE SET NULL  
  ON UPDATE CASCADE)
```

Oracle does not support "on update"

# In-Class Exercise (no need to hand it in)




Translate the ER diagram to relational. Underline key attributes, and make FKs bold (or circle the FK if you are doing this by hand).

Back at 3:01

Work with the people around you!

# Clicker Question



Translate the ER diagram to relational.


Which of the following appears in your relational schema:

- A.  $AR(\underline{a}, \mathbf{b}, d)$
- B.  $BS(\underline{b}, \mathbf{c}, e)$
- C.  $S(\underline{b}, \underline{c})$
- D. All of these
- E. None of these

Primary keys are underlined. Foreign keys are bolded.



# Clicker Question



Translate the ER diagram to relational.

Which of the following appears in your relational schema:

A.  $AR(\underline{a}, \underline{b}, d)$

B.  $BS(\underline{b}, \underline{c}, e)$

C.  $S(\underline{b}, \underline{c})$

D. All of these

E. None of these


~~A  
B  
C  
R  
S~~

AR  
BS  
C


→

$AR(\underline{a}, \underline{b}, d)$   
 $BS(\underline{b}, \underline{c}, e)$   
 $C(\underline{c}, f)$

# Relationship Sets with Key Constraints (one to one case)



# Relationship Sets with Key Constraints (one to one case)




Which schema below is a reasonable translation from ER to relations? (bolded attributes are foreign keys)

- A. Country(coName, **caName**)
- B. Country(name), Capital(name)
- C. Capital (caName, **coName**)
- D. Both A and C**
- E. All of A, B, and C



# Relationship Sets with Key Constraints (one to one case)



Let's assume we went with Country(coName, caName). Do we need a separate relation for Capital?

- A. Yes
- B. No
- C. It depends

In this case there are no additional attributes in Capital, and both have total participation so no. But if there were extra Capital attributes or there was not total participation, maybe yes

The goal on the exam is to ask you questions where the answer is clear!