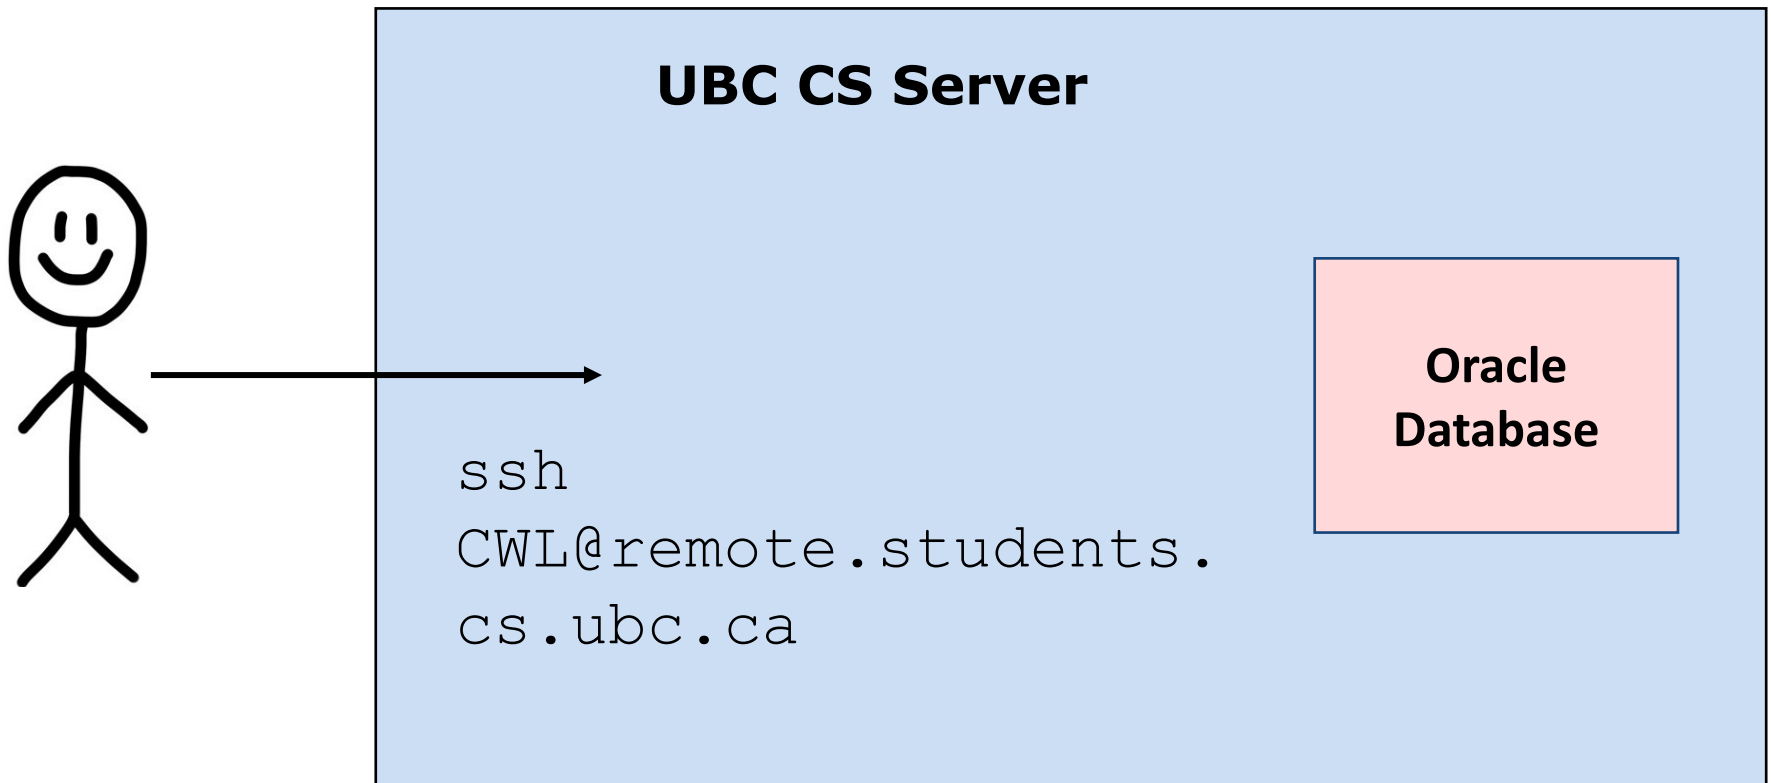




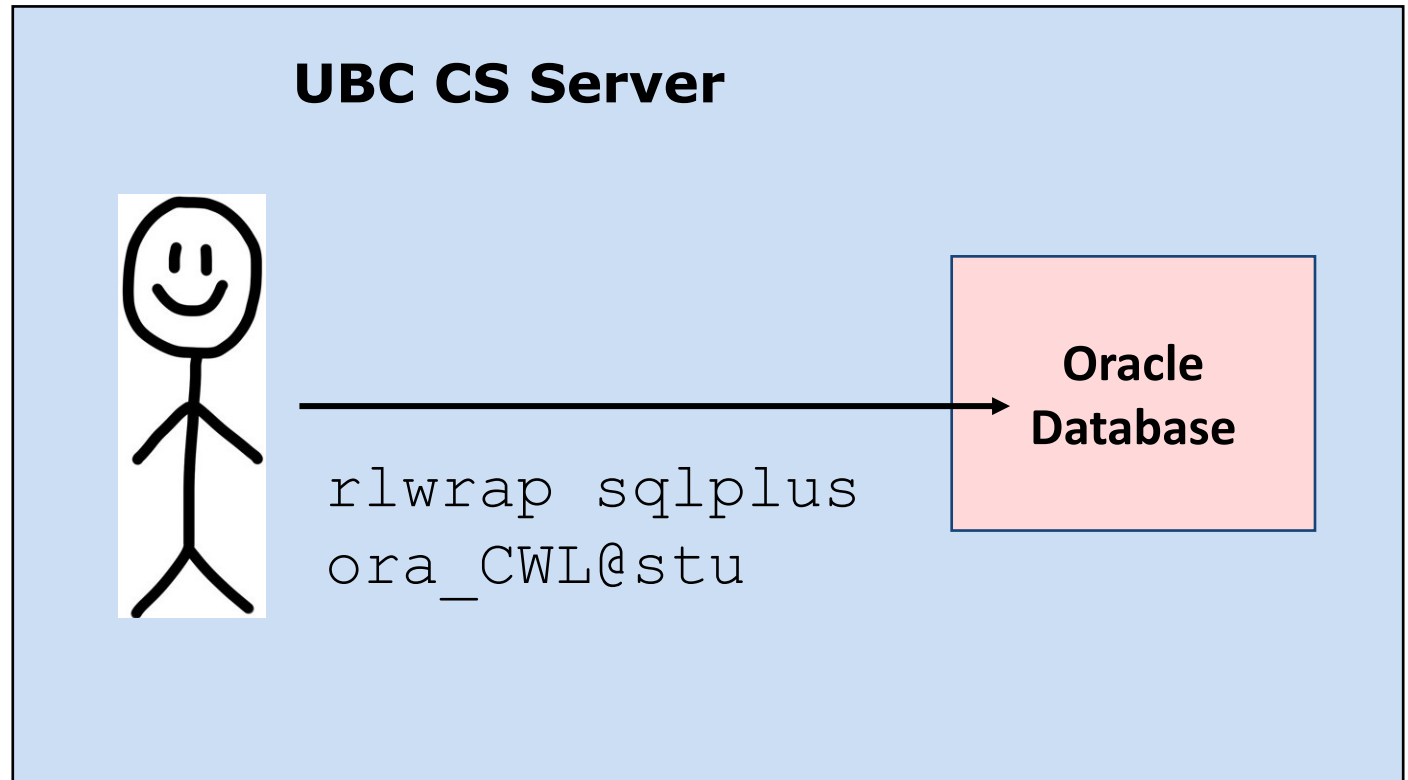
Administrative Notes – February 7, 2023

- Feb 17: Assignment 3 (pairs if you want) due
 - Make sure you sign up in pairs beforehand as we need time to configure Canvas to allow both partners to view the submission and subsequent feedback
 - Everyone who has signed up as of last night has been put into a group on Canvas
- Feb 20 – 24: Reading Break! (yay!)
 - No lectures, tutorials, or office hours during this week

Logging into SQL Plus

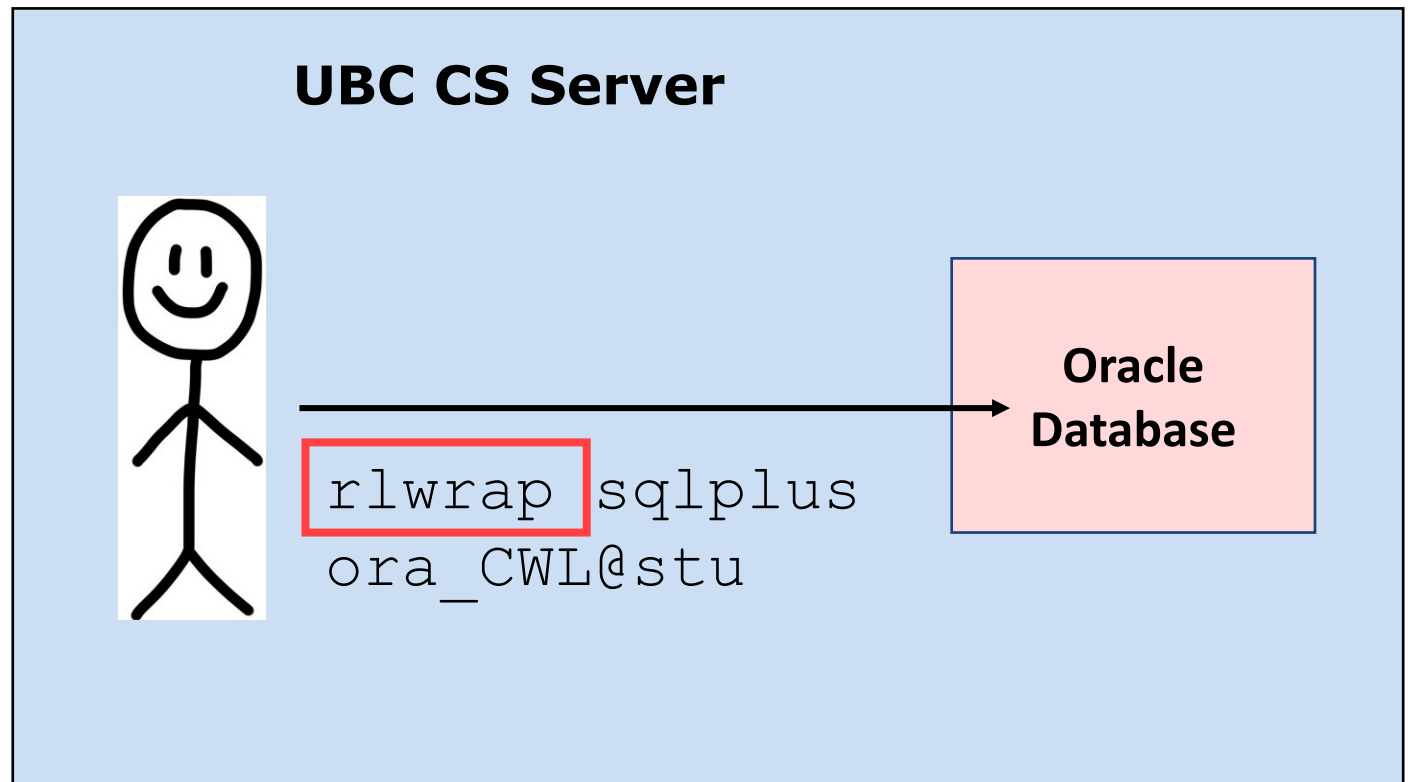


Logging into SQL Plus

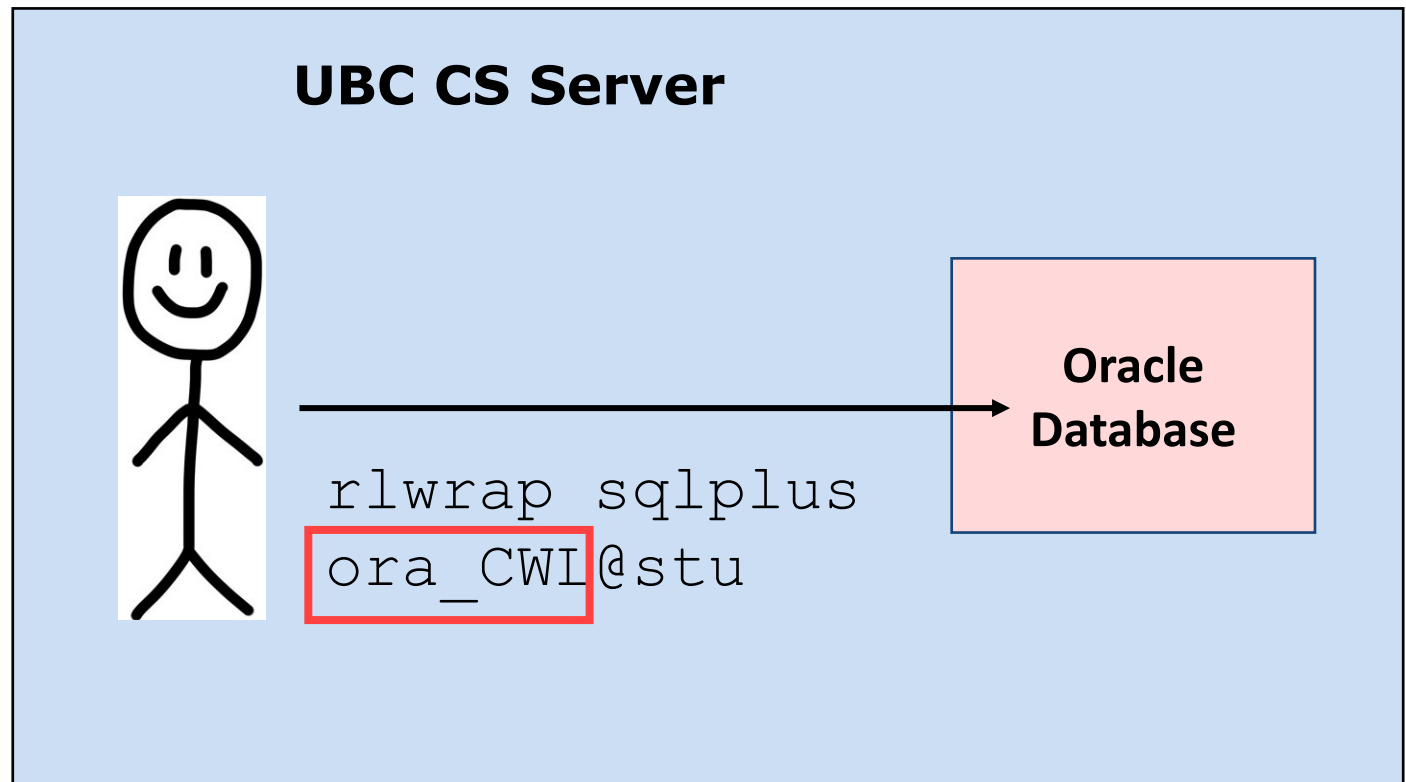


Password is a<student number>
like a12345678.

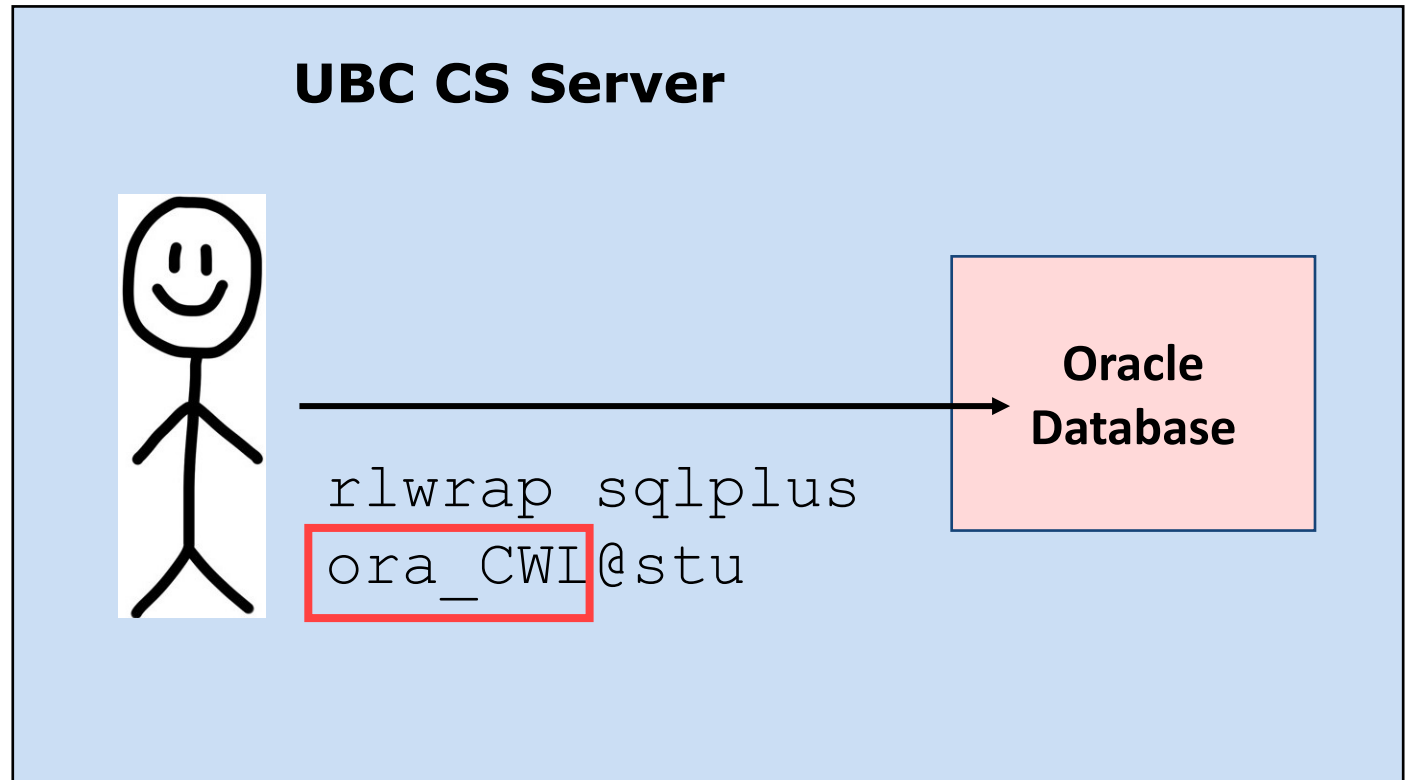
Logging into SQL Plus




Logging into SQL Plus



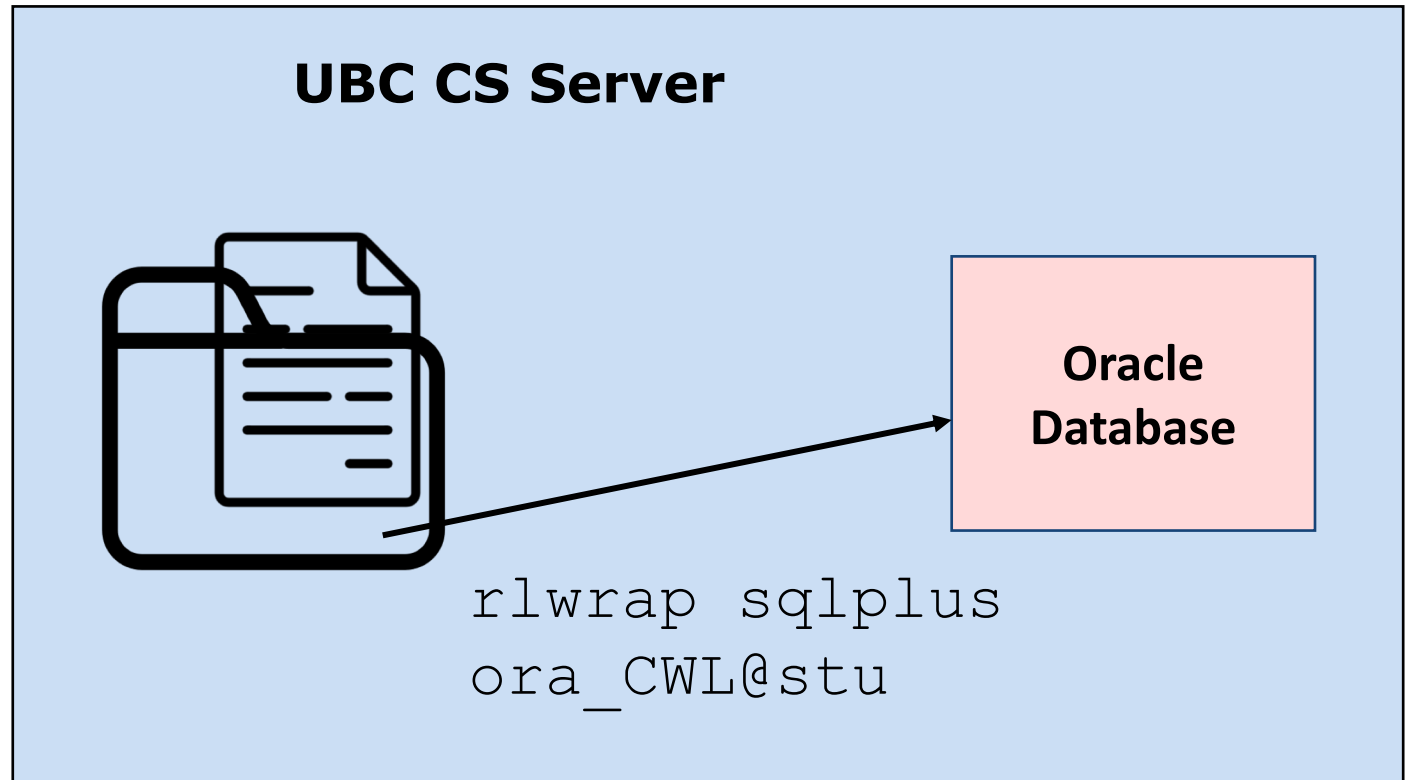
If you want to run a file of SQL statements
in Oracle...



If you want to run a file of SQL statements
in Oracle...



If you want to run a file of SQL statements
in Oracle...





Once you have logged into SQLPlus...

Oracle Database

```
start <filename>.sql;
```

Tutorial will be focused on SQLPlus for the next two weeks. Go to tutorial to get help if you are still stuck!



Now where were we...

- SELECT, FROM, WHERE

SELECT	<i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

- We know how to work across multiple relations
- We know how to alias a relation (super important when relations have the same attribute name)
- We know how to rename any attributes in our result

Conceptual Procedural Evaluation Strategy

1. Compute the cross-product of *relation-list*.
2. Discard resulting tuples if they fail *qualifications*.
3. Delete attributes that are not in *target-list*.
4. If **DISTINCT** is specified, eliminate duplicate rows.

SELECT	<i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>



Example of Conceptual Procedural Evaluation

```
SELECT Name
FROM MovieStar M, StarsIn S
WHERE S.StarID = M.StarID AND
      MovieID = 276
```

StarID	Name	Gender	MovieID	StarID	Character
1273	Nathalie Portman	Female	272	1269	Leigh Anne Touhy
1273	Nathalie Portman	Female	273	1270	Mary
1273	Nathalie Portman	Female	274	1271	King George VI
1273	Nathalie Portman	Female	276	1273	Nina Sayers
...



New Students Example

- Class(name,meets_at,room,fid)
- Student(snum,sname,major,standing,age)
- Enrolled(snum,cname)
- Faculty(fid,fname,deptid)



Class Table

Name	Meets_at	Room	FID
Data Structures	MWF 10	R128	489456522
Database Systems	MWF 12:30-1:45	1320 DCL	142519864
Operating System Design	TuTh 12-1:20	20 AVW	489456522
Archaeology of the Incas	MWF 3-4:15	R128	248965255
Aviation Accident Investigation	TuTh 1-2:50	Q3	011564812
Air Quality Engineering	TuTh 10:30-11:45	R15	011564812
Introductory Latin	MWF 3-4:15	R12	248965255
American Political Parties	TuTh 2-3:15	20 AVW	619023588
Social Cognition	Tu 6:30-8:40	R15	159542516
Perception	MTuWTh 3	Q3	489221823
Multivariate Analysis	TuTh 2-3:15	R15	090873519
Patent Law	F 1-2:50	R128	090873519
Urban Economics	MWF 11	20 AVW	489221823
Organic Chemistry	TuTh 12:30-1:45	R12	489221823
Marketing Research	MW 10-11:15	1320 DCL	489221823
Seminar in American Art	M 4	R15	489221823
Orbital Mechanics	MWF 8 1320	DCL	011564812
Dairy Herd Management	TuTh 12:30-1:45	R128	356187925
Communication Networks	MW 9:30-10:45	20 AVW	141582651
Optical Electronics	TuTh 12:30-1:45	R15	254099823
Introduction to Math	TuTh 8-9:30	R128	489221823



Student Table

SNUM	SNAME	MAJOR	ST	AGE
51135593	Maria White	English	SR	21
60839453	Charles Harris	Architecture	SR	22
99354543	Susan Martin	Law	JR	20
112348546	Joseph Thompson	Computer Science	SO	19
115987938	Christopher Garcia	Computer Science	JR	20
132977562	Angela Martinez	History	SR	20
269734834	Thomas Robinson	Psychology	SO	18
280158572	Margaret Clark	Animal Science	FR	18
301221823	Juan Rodriguez	Psychology	JR	20
318548912	Dorothy Lewis	Finance	FR	18
320874981	Daniel Lee	Electrical Engineering	FR	17
322654189	Lisa Walker	Computer Science	SO	17
348121549	Paul Hall	Computer Science	JR	18
351565322	Nancy Allen	Accounting	JR	19
451519864	Mark Young	Finance	FR	18
455798411	Luis Hernandez	Electrical Engineering	FR	17
462156489	Donald King	Mechanical Engineering	SO	19
550156548	George Wright	Education	SR	21
552455318	Ana Lopez	Computer Engineering	SR	19
556784565	Kenneth Hill	Civil Engineering	SR	21
567354612	Karen Scott	Computer Engineering	FR	18
573284895	Steven Green	Kinesiology	SO	19
574489456	Betty Adams	Economics	JR	20
578875478	Edward Baker	Veterinary Medicine	SR	21



Enrolled Table

SNUM	CNAME
112348546	Database Systems
115987938	Database Systems
348121549	Database Systems
322654189	Database Systems
552455318	Database Systems
455798411	Operating System Design
552455318	Operating System Design
567354612	Operating System Design
112348546	Operating System Design
115987938	Operating System Design
322654189	Operating System Design
567354612	Data Structures
552455318	Communication Networks
455798411	Optical Electronics
455798411	Organic Chemistry
301221823	Perception
301221823	Social Cognition
301221823	American Political Parties
556784565	Air Quality Engineering
99354543	Patent Law
574489456	Urban Economics



Faculty Table

FID	FNAME	DEPTID
142519864	I. Teach	20
242518965	James Smith	68
141582651	Mary Johnson	20
011564812	John Williams	68
254099823	Patricia Jones	68
356187925	Robert Brown	12
489456522	Linda Davis	20
287321212	Michael Miller	12
248965255	Barbara Wilson	12
159542516	William Moore	33
090873519	Elizabeth Taylor	11
486512566	David Anderson	20
619023588	Jennifer Thomas	11
489221823	Richard Jackson	33
548977562	Ulysses Teach	20



Running Examples

Movie(MovieID, Title, Year)

StarsIn(MovieID, StarID, Character)

MovieStar(StarID, Name, Gender)

Student(snum, sname, major, standing, age)

Class(name, meets_at, room, fid)

Enrolled(snum, cname)

Faculty(fid, fname, deptid)



What kinds of queries can you answer so far?

Find the student ids of those who have taken a course named “Database Systems”.

Student(snum,sname,major,standing,age)

Class(name,meets_at,room,fid)

Enrolled(snum,cname)

Faculty(fid,fname,deptid)



What kinds of queries can you answer so far?

Find the student ids of those who have taken a course named “Database Systems”.

```
SELECT snum  
FROM enrolled e  
WHERE cname = 'Database Systems'
```

Do we need distinct? A. Yes. **B. No**

Student(snum,sname,major,standing,age)

Class(name,meets_at,room,fid)

Enrolled(snum,cname)

Faculty(fid,fname,deptid)



What kinds of queries can you answer so far?

Find the names of all classes taught by Elizabeth Taylor.

Student(snum,sname,major,standing,age)

Class(name,meets_at,room,fid)

Enrolled(snum,cname)

Faculty(fid,fname,deptid)



What kinds of queries can you answer so far?

Find the names of all classes taught by Elizabeth Taylor.

```
SELECT name
FROM Faculty f, class c
WHERE f.fid = c.fid and
      fname = 'Elizabeth Taylor'
```

Do we need distinct? A. Yes.

B. No

Student(snum,sname,major,standing,age)

Class(name,meets_at,room,fid)

Enrolled(snum,cname)

Faculty(fid,fname,deptid)



What kinds of queries can you answer so far?

Find the departments that have at least one faculty member

Student(snum,sname,major,standing,age)

Class(name,meets_at,room,fid)

Enrolled(snum,cname)

Faculty(fid,fname,deptid)



What kinds of queries can you answer so far?

Find the departments that have at least one faculty member

```
SELECT DISTINCT deptid  
FROM faculty
```

Do we need distinct? **A. Yes.** B. No

Student(snum,sname,major,standing,age)

Class(name,meets_at,room,fid)

Enrolled(snum,cname)

Faculty(fid,fname,deptid)



What kinds of queries can you answer so far?

Find the departments that have more than one faculty member (express not equal by “<>”)

```
SELECT DISTINCT f1.deptid
FROM faculty f1, faculty f2
WHERE f1.fid <> f2.fid AND
      f1.deptid = f2.deptid
```

That is why
renaming is
important.

f1

<u>fid</u>	fname	Deptid
90873519	Elizabeth Taylor	11
61902358	Jennifer Thomas	11
8		
...

f2

<u>fid</u>	fname	Deptid
90873519	Elizabeth Taylor	11
61902358	Jennifer Thomas	11
8		
...

Do you need DISTINCT?



String Comparisons

What are the student ids of those who have taken a course with “System” in the name?



A string walks into a bar...

```
SELECT DISTINCT snum  
FROM   enrolled  
WHERE  cname LIKE '%System%'
```

- **LIKE** is used for string matching:
 - **'_'** stands for any one character and
 - **'%'** stands for 0 or more arbitrary characters.
- SQL supports string operations such as
 - concatenation (using "||")
 - converting from upper to lower case (and vice versa)
 - finding string length, extracting substrings, etc.



A string walks into a bar...

```
SELECT DISTINCT snum  
FROM   enrolled  
WHERE  cname LIKE '%System%'
```

Do we need DISTINCT?

A. Yes

B. No



Ordering of Tuples

List in alphabetic order **the names of actors who were in a movie in 1939.**

```
SELECT DISTINCT name
FROM Movie, StarsIn, MovieStar
WHERE Movie.MovieID = StarsIn.MovieID AND
      StarsIn.StarID = MovieStar.StarID AND
      year = 1939
```



Ordering of Tuples

List in **alphabetic order** the names of actors who were in a movie in 1939.

```
SELECT DISTINCT name
FROM Movie, StarsIn, MovieStar
WHERE Movie.MovieID = StarsIn.MovieID AND
      StarsIn.StarID = MovieStar.StarID AND
      year = 1939
ORDER BY name
```



Ordering of Tuples

Order is specified by:

- **desc** for descending order
- **asc** for ascending order (default)
- E.g. **order by *Name* desc**
- You can order within order: for example, ... “ORDER BY Year, Name” would first order by Year, then Name within years



Clicker Question: Sorting

Relation R has schema R(a,b,c). In the result of the query

```
SELECT a, b, c
FROM R
ORDER BY c DESC, b ASC;
```

What condition must a tuple t satisfy so that t **necessarily precedes** the tuple (5,5,5)? Identify one such tuple from the list below.

- A. (3,6,3)
- B. (1,5,5)
- C. (5,5,6)
- D. All of the above
- E. None of the above



Clicker Question: Sorting

Relation R has schema R(a,b,c). In the result of the query

```
SELECT a, b, c
FROM R
ORDER BY c DESC, b ASC;
```

What condition must a tuple t satisfy so that t **necessarily precedes** the tuple (5,5,5)? Identify one such tuple from the list below.

A. (3,6,3) 3 < 5

B. (1,5,5) Not specified

C. (5,5,6) Right

D. All of the above

E. None of the above

clickerorder.sql and
clickerorder2.sql produce
different ordering for
7,5,5 vs. 1,5,5



Set Operations

- **union, intersect, and except** correspond to the relational algebra operations \cup , \cap , $-$.
- Each automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions:

union all, intersect all and except all

- **A Union B (OR)** – combine the results from A and B
- **A Intersect B (AND)** – only keep results that appear in both A and B
- **A Except B** – only keep results found in A and not B



Set Operations

- Suppose a tuple occurs m times in r and n times in s , then, it occurs:
 - $m + n$ times in r **union all** s
 - $\min(m, n)$ times in r **intersect all** s
 - $\max(0, m - n)$ times in r **except all** s



Find IDs of MovieStars who've been in a movie in 1944 or 1974

- **UNION:** Can union any two *union-compatible* sets of tuples (i.e., the result of SQL queries).

```
Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Character)
MovieStar(StarID, Name, Gender)
```



Find IDs of MovieStars who've been in a movie in 1944 or 1974

```
Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Character)
MovieStar(StarID, Name, Gender)
```

- **UNION:** Can union any two *union-compatible* sets of tuples (i.e., the result of SQL queries).

```
SELECT StarID
FROM Movie M, StarsIn S
WHERE M.MovieID=S.MovieID AND
( year = 1944 OR year = 1974)
```

- The two queries though quite similar return different results, why?
 - Use UNION ALL to get the same answer

```
SELECT StarID
FROM Movie M, StarsIn S
WHERE M.MovieID = S.MovieID AND
year = 1944
UNION
SELECT StarID
FROM Movie M, StarsIn S
WHERE M.MovieID = S.MovieID AND
year = 1974
```



Set Operations: Intersect

Example: Find IDs of stars who have been in a movie in 1944 and 1974.

```
Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Character)
MovieStar(StarID, Name, Gender)
```

- **INTERSECT:** Can be used to compute the intersection of any two *union-compatible* sets of tuples.
- In SQL/92, but some systems don't support it.



Set Operations: Intersect

Example: Find IDs of stars who have been in a movie in 1944 and 1974.

```
Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Character)
MovieStar(StarID, Name, Gender)
```

- **INTERSECT**: Can be used to compute the intersection of any two *union-compatible* sets of tuples.
- In SQL/92, but some systems don't support it.

```
SELECT StarID
FROM    Movie M, StarsIn S
WHERE   M.MovieID = S.MovieID AND
year = 1944
INTERSECT
SELECT StarID
FROM    Movie M, StarsIn S
WHERE   M.MovieID = S.MovieID AND
year = 1974
```

Oracle does
MYSQL doesn't



Rewriting INTERSECT with Joins

Example: Find IDs of stars who have been in a movie in 1944 and 1974 without using **INTERSECT**.

```
Movie(MovieID, Title, Year)
StarsIn(MovieID, StarID, Character)
MovieStar(StarID, Name, Gender)
```




Rewriting INTERSECT with Joins

Example: Find IDs of stars who have been in a movie in 1944 and 1974 without using **INTERSECT**.

```
SELECT distinct S1.StarID
FROM    Movie M1, StarsIn S1,
        Movie M2, StarsIn S2
WHERE
        M1.MovieID = S1.MovieID AND M1.year = 1944 AND
        M2.MovieID = S2.MovieID AND M2.year = 1974 AND
        S2.StarID = S1.StarID
```



Set Operations: EXCEPT

Find the sids of all students who took Operating System Design but did not take Database Systems

Student(snum,sname,major,standing,age)

Class(name,meets_at,room,fid)

Enrolled(snum,cname)

Faculty(fid,fname,deptid)



Set Operations: EXCEPT

Find the sids of all students who took Operating System Design but did not take Database Systems

```
SELECT snum
FROM enrolled e
WHERE cname = 'Operating System Design'
EXCEPT ← Oracle uses MINUS rather than EXCEPT
SELECT snum
FROM enrolled e
WHERE cname = 'Database Systems'
```

Can we do it in a different way?
(We'll come back to this)



But what about...

Select the IDs of all students who have not taken “Operating System Design”

- One way to do is to find all students that have taken a course.
- MINUS those who have taken “Operating System Design”

```
SELECT snum
FROM enrolled e
EXCEPT ← Oracle uses MINUS rather than EXCEPT
SELECT snum
FROM enrolled e
WHERE cname = 'Operating System
              Design'
```



Motivating Example for Nested Queries

Find ids and names of female stars who have been in movie with ID 28:



Motivating Example for Nested Queries

Find ids and names of female stars who have been in movie with ID 28:

```
SELECT M.StarID, name
FROM MovieStar M, StarsIn S
WHERE M.StarID = S.starID AND S.MovieID = 28
      AND gender = 'female';
```



Motivating Example for Nested Queries

Find ids and names of female stars who have been in movie with ID 28:

```
SELECT M.StarID, name
FROM MovieStar M, StarsIn S
WHERE M.StarID = S.starID AND S.MovieID = 28
      AND gender = 'female';
```

Find ids and names of female stars who have not been in movie w/ ID 28 w/o using EXCEPT/MINUS:

- Would the following be correct?

```
SELECT M.StarID, name
FROM MovieStar M, StarsIn S
WHERE M.StarID = S.starID AND S.MovieID <> 28
      AND gender = 'female';
```



Motivating Example for Nested Queries

MovieStar

<u>StarID</u>	Name	Gender
1	Jessica Wong	Female
2	Jia Lu	Male
3	Carol Huang	Female

StarsIn

<u>MovieID</u>	StarID	Character
28	1	A tree
28	3	Background
29	3	A tired grad student

Does “WHERE
StarsIn.MovieID
<> 28” correctly
remove Carol from
the results?



Nested Queries

- A very powerful feature of SQL:

```
SELECT  A1, A2, ..., An
FROM    R1, R2, ..., Rm
WHERE   condition
```

- A nested query is a query that has another query embedded with it.
 - A **SELECT, FROM, WHERE, or HAVING** clause can itself contain an SQL query!
 - Being part of the **WHERE** clause is the most common



Nested Queries (IN/Not IN)

Find ids and names of stars who have been in movie with ID 28:



Nested Queries (IN/Not IN)

Find ids and names of stars who have been in movie with ID 28:

```
SELECT M.StarID, M.Name
FROM MovieStar M
WHERE M.Gender = 'female' AND
      M.StarID IN (SELECT S.StarID
                   FROM StarsIn S
                   WHERE MovieID = 28)
```

There's also NOT IN

- To find stars who have *not* been in movie 28, use **NOT IN**.
- To understand nested query semantics, think of a nested loops evaluation:
 - *For each MovieStar tuple, check the qualification by computing the subquery.*