



Administrative Notes – February 14, 2023

- Feb 17: Assignment 3 (pairs if you want) due
 - Make sure you sign up in pairs beforehand as we need time to configure Canvas to allow both partners to view the submission and subsequent feedback
 - Everyone who has signed up as of last night has been put into a group on Canvas
- Feb 17: Assignment 4 will be released
- Feb 20 – 24: Reading Break! (yay!)
 - No lectures, tutorials, or office hours during this week



Review: Examples for Division

A

| sno | pno |
|-----|-----|
| s1 | p1 |
| s1 | p2 |
| s1 | p3 |
| s1 | p4 |
| s2 | p1 |
| s2 | p2 |
| s3 | p2 |
| s4 | p2 |
| s4 | p4 |

B1

| pno |
|-----|
| p2 |

B2

| pno |
|-----|
| p2 |
| p4 |

B3

| pno |
|-----|
| p1 |
| p2 |
| p4 |

A/B1

| sno |
|-----|
| s1 |
| s2 |
| s3 |
| s4 |

A/B2

| sno |
|-----|
| s1 |
| s4 |

A/B3

| sno |
|-----|
| s1 |

Review: Division in SQL

Find students who've
taken **all** classes.

```
SELECT  sname
FROM    Student S
WHERE   NOT EXISTS
        ((SELECT  C.name
          FROM    Class C)
         EXCEPT
         (SELECT  E.cname
          FROM    Enrolled E
          WHERE   E.snum=S.snum) )
```

(method 1)

The hard way (without EXCEPT):

```
SELECT  sname
FROM    Student S
WHERE   NOT EXISTS (
    SELECT  C.name
    FROM    Class C
    WHERE   NOT EXISTS (SELECT  E.snum
                        FROM    Enrolled E
                        WHERE   C.name= E.cname AND
                                E.snum=S.snum) )
```

(method 2)



In-Class Exercise (SQL 3)

- Canvas → Modules → In Class Exercises
- You can work on it with other people around you. If you work with others, you must **write their names on your submission to acknowledge the collaboration.**
 - Everyone must submit to Canvas
- Reminder: no late submissions accepted



Clicker Question

I am ready to cover the in-class exercise.

- A. Yes
- B. No
- C. I need two more minutes
- D. I need five more minutes



Aggregate Operators

- These functions operate on the multiset of values of a column of a relation, and return a value

AVG: average value

MIN: minimum value

MAX: maximum value

SUM: sum of values

COUNT: number of values

```
SELECT count(s.snum)
FROM enrolled e, Student S
WHERE e.snum = s.snum
```

- The following versions eliminate duplicates before applying the operation to attribute A:

COUNT (DISTINCT A)

SUM (DISTINCT A)

AVG (DISTINCT A)

```
SELECT count(distinct s.snum)
FROM enrolled e, Student S
WHERE e.snum = s.snum
```



Aggregate Operators: Examples

students

```
SELECT  COUNT (*)  
FROM    Student
```

Find name and age
of the oldest
student(s).

```
SELECT  sname, age  
FROM    Student S  
WHERE   S.age = (SELECT  MAX(S2.age)  
                FROM    Student S2)
```

Finding average
age
of SR students.

```
SELECT  AVG (age)  
FROM    Student  
WHERE   standing='SR'
```



Aggregation Examples

Find the minimum student age.

How many students have taken a class with
“Database” in the title.

Student(snum,sname,major,standing,age)

Class(name,meets_at,room,fid)

Enrolled(snum,cname)

Faculty(fid,fname,deptid)



Aggregation Examples

Find the minimum student age.

```
SELECT MIN(age)
FROM student;
```

How many students have taken a class with
“Database” in the title.

```
SELECT COUNT(DISTINCT snum)
FROM enrolled
WHERE cname LIKE '%Database%'
```



GROUP BY and HAVING

- Divide tuples into groups and apply aggregate operations to each group.
- Example: *Find the age of the youngest student for each major.*

For $i = \text{'Computer Science'},$
 $\text{'Civil Engineering'} \dots$

```
SELECT  MIN (age)
FROM    Student
WHERE   major = i
```

Problem:

We don't know how many majors exist, not to mention this is not good practice



Grouping Examples

Find the age of the youngest student who is at least 19, for each major.

```
SELECT major, MIN(age)
FROM Student
WHERE age >= 19
GROUP BY major
```

| Snum | Major | Age |
|-----------|-------------------|-----|
| 115987938 | Computer Science | 20 |
| 112348546 | Computer Science | 19 |
| 280158572 | Animal Science | 18 |
| 351565322 | Accounting | 19 |
| 556784565 | Civil Engineering | 21 |
| ... | ... | ... |

No Animal Science

| Major | Age |
|-------------------|-----|
| Computer Science | 19 |
| Accounting | 19 |
| Civil Engineering | 21 |
| ... | ... |



Grouping Examples with Having

Find the age of the youngest student who is at least 19, for each major with at least 2 such students.

```
SELECT major, MIN(age)
FROM Student
WHERE age >= 19
GROUP BY major
HAVING COUNT(*) > 1
```

| Snum | Major | Age |
|-----------|-------------------|-----|
| 115987938 | Computer Science | 20 |
| 112348546 | Computer Science | 19 |
| 280158572 | Animal Science | 18 |
| 351565322 | Accounting | 19 |
| 556784565 | Civil Engineering | 21 |
| ... | ... | ... |



Grouping Examples with Having

Find the age of the youngest student who is at least 19, for each major with at least 2 such students.

```
SELECT major, MIN(age)
FROM Student
WHERE age >= 19
GROUP BY major
HAVING COUNT(*) > 1
```

| Snum | Major | Age |
|-----------|-------------------|-----|
| 115987938 | Computer Science | 20 |
| 112348546 | Computer Science | 19 |
| 280158572 | Animal Science | 18 |
| 351565322 | Accounting | 19 |
| 556784565 | Civil Engineering | 21 |
| ... | ... | ... |

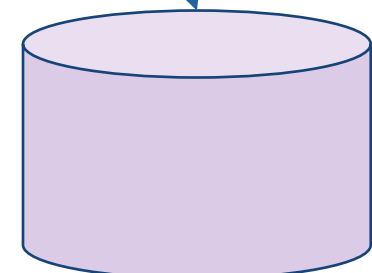
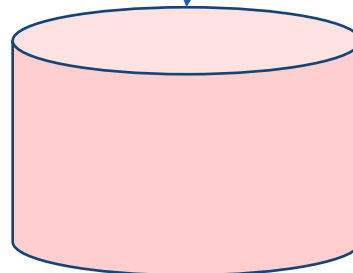
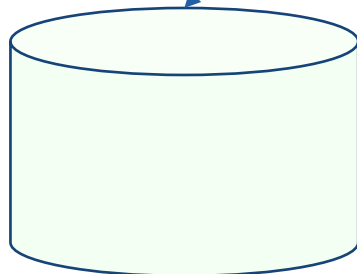


Grouping Examples with Having

Find the age of the youngest student who is at least 19, for each major with at least 2 such students.

```
SELECT major, MIN(age)
FROM Student
WHERE age >= 19
GROUP BY major
HAVING COUNT(*) > 1
```

| Snum | Major | Age |
|-----------|-------------------|-----|
| 115987938 | Computer Science | 20 |
| 112348546 | Computer Science | 19 |
| 280158572 | Animal Science | 18 |
| 351565322 | Accounting | 19 |
| 556784565 | Civil Engineering | 21 |
| ... | ... | ... |



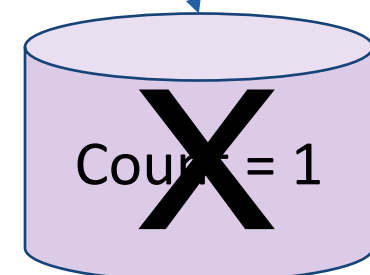
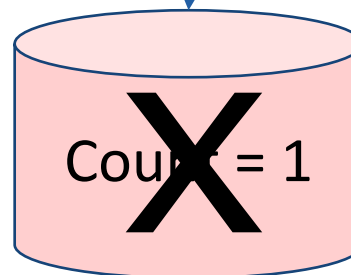
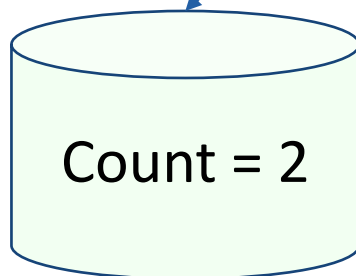


Grouping Examples with Having

Find the age of the youngest student who is at least 19, for each major with at least 2 such students.

```
SELECT major, MIN(age)
FROM Student
WHERE age >= 19
GROUP BY major
HAVING COUNT(*) > 1
```

| Snum | Major | Age |
|-----------|-------------------|-----|
| 115987938 | Computer Science | 20 |
| 112348546 | Computer Science | 19 |
| 280158572 | Animal Science | 18 |
| 351565322 | Accounting | 19 |
| 556784565 | Civil Engineering | 21 |
| ... | ... | ... |





Grouping Examples with Having

Find the age of the youngest student who is at least 19, for each major with at least 2 such students.

```
SELECT major, MIN(age)
FROM Student
WHERE age >= 19
GROUP BY major
HAVING COUNT(*) > 1
```

| Snum | Major | Age |
|-----------|-------------------|-----|
| 115987938 | Computer Science | 20 |
| 112348546 | Computer Science | 19 |
| 280158572 | Animal Science | 18 |
| 351565322 | Accounting | 19 |
| 556784565 | Civil Engineering | 21 |
| ... | ... | ... |

| Major | Age |
|------------------|-----|
| Computer Science | 19 |



And there are rules

Find the age of the youngest student who is at least 19, for each major with at least 2 such students

```
SELECT    major, MIN(age)
FROM      Student
WHERE     age >= 19
GROUP BY  major
HAVING    COUNT(*) > 1
```

- Would it make sense if I select age instead of MIN(age)?
- *Would it make sense if I select snum to be returned?*
- *Would it make sense if I select major to be returned?*



And there are rules

Find the age of the youngest student who is at least 19, for each major with at least 2 such students.

```
SELECT major, MIN(age)
FROM Student
WHERE age >= 19
GROUP BY major
HAVING COUNT(*) > 1
```

- Would it make sense if I select `age` instead of `MIN(age)` ?
- Would it make sense if I select `snum` to be returned?
- Would it make sense if I select `major` to be returned?

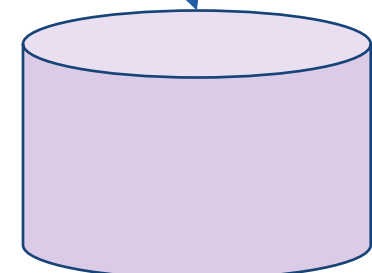
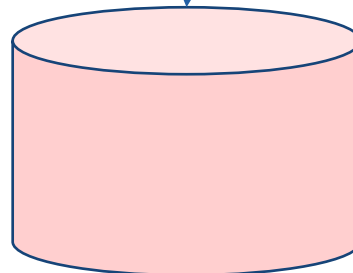
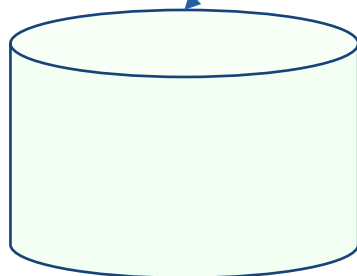


And there are rules

Find the age of the youngest student who is at least 19, for each major with at least 2 such students.

```
SELECT major, MIN(age)
FROM Student
WHERE age >= 19
GROUP BY major
HAVING COUNT(*) > 1
```

| Snum | Major | Age |
|-----------|-------------------|-----|
| 115987938 | Computer Science | 20 |
| 112348546 | Computer Science | 19 |
| 280158572 | Animal Science | 18 |
| 351565322 | Accounting | 19 |
| 556784565 | Civil Engineering | 21 |
| ... | ... | ... |





GROUP BY and HAVING (cont)

```
SELECT [DISTINCT]  target-list
FROM relation-list
WHERE qualification
GROUP BY grouping-list
HAVING group-qualification
ORDER BY target-list
```

- The *target-list* contains
 - (i) attribute names
 - (ii) terms with aggregate operations (e.g., MIN (S . age)).
- Attributes in (i) must also be in *grouping-list*.
 - each answer tuple corresponds to a *group*,
 - *group* = a set of tuples with same value for all attributes in *grouping-list*
 - selected attributes must have a single value per group.
- Attributes in *group-qualification* are either in *grouping-list* or are arguments to an aggregate operator.

Conceptual Evaluation of a Query

1. compute the cross-product of *relation-list* **from**
2. keep only tuples that satisfy *qualification* **where**
3. partition the remaining tuples into groups by the value of attributes in *grouping-list* **group by**
4. keep only the groups that satisfy *group-qualification*
(expressions in *group-qualification* must have a *single value per group!*) **having**
5. delete fields that are not in *target-list* **select**
6. generate one answer tuple per qualifying group.



Out of Order SQL Clauses

6 rows selected.


```
SQL> ;
  1 SELECT a1.x, a2.y, COUNT(*)
  2 FROM Arc a1, Arc a2
  3 WHERE a1.y = a2.x
  4* GROUP BY a1.x, a2.y
SQL>
SQL>
SQL>
SQL> SELECT a1.x, a2.y, COUNT(*)
  2 FROM Arc a1, Arc a2
  3 GROUP BY a1.x, a2.y
  4 WHERE a1.y = a2.x;
WHERE a1.y = a2.x
*
ERROR at line 4:
ORA-00933: SQL command not properly ended
```

```
SQL> SELECT b, count(c)
  2 FROM R
  3 WHERE B > 5
  4 GROUP BY B
  5 ORDER BY count(c);
```

| B | COUNT(C) |
|---|----------|
| 6 | 1 |

```
SQL> SELECT b, count(c)
  2 FROM R
  3 WHERE B > 5
  4 ORDER BY count(c)
  5 GROUP BY B;
GROUP BY B
*
ERROR at line 5:
ORA-00933: SQL command not properly ended
```

GROUP BY and HAVING



```
Student(snum,sname,major,standing,age)
Class(name,meets_at,room,fid)
Enrolled(snum,cname)
Faculty(fid,fname,deptid)
```

Example 1: For each class, find the age of the youngest student who has enrolled in this class.

```
SELECT cname, MIN(age)
FROM Student S, Enrolled E
WHERE S.snum= E.snum
GROUP BY cname
```

Example 2: For each course with more than 1 enrollment, find the age of the youngest student who has taken this class.

```
SELECT cname, MIN(age)
FROM Student S, Enrolled E
WHERE S.snum = E.snum
GROUP BY cname
HAVING COUNT(*) > 1
```

← per group qualification!



Clicker Question: Grouping

clickergrouping.sql

Compute the result of the query:

```
SELECT a1.x, a2.y, COUNT(*)  
FROM Arc a1, Arc a2  
WHERE a1.y = a2.x  
GROUP BY a1.x, a2.y
```

Which of the following is in the result?

A. (1,3,2)

B. (4,2,6)

C. (4,3,1)

D. All of the above

E. None of the above

| x | y | COUNT(*) |
|---|---|----------|
| 1 | 3 | 2 |
| 2 | 4 | 2 |
| 3 | 1 | 6 |
| 3 | 2 | 2 |
| 4 | 2 | 6 |
| 4 | 3 | 1 |

Arc

| x | y |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 2 | 3 |
| 3 | 4 |
| 3 | 4 |
| 4 | 1 |
| 4 | 1 |
| 4 | 1 |
| 4 | 2 |



Clicker Question: Grouping

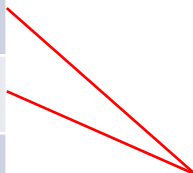
```
SELECT a1.x, a2.y, COUNT(*)  
FROM Arc a1, Arc a2  
WHERE a1.y = a2.x  
GROUP BY a1.x, a2.y
```

Arc a1

| x | y |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 2 | 3 |
| 3 | 4 |
| 3 | 4 |
| 4 | 1 |
| 4 | 1 |
| 4 | 1 |
| 4 | 2 |

Arc a2

| x | y |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 2 | 3 |
| 3 | 4 |
| 3 | 4 |
| 4 | 1 |
| 4 | 1 |
| 4 | 1 |
| 4 | 2 |



Which of the following is in the result?

- A. (1,3,2)
- B. (4,2,6)
- C. (4,3,1)
- D. All of the above
- E. None of the above



Clicker Question: Grouping

```
SELECT a1.x, a2.y, COUNT(*)  
FROM Arc a1, Arc a2  
WHERE a1.y = a2.x  
GROUP BY a1.x, a2.y
```

Arc a1

| x | y |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 2 | 3 |
| 3 | 4 |
| 3 | 4 |
| 4 | 1 |
| 4 | 1 |
| 4 | 1 |
| 4 | 2 |

Arc a2

| x | y |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 2 | 3 |
| 3 | 4 |
| 3 | 4 |
| 4 | 1 |
| 4 | 1 |
| 4 | 1 |
| 4 | 2 |



Which of the following is in the result?

- A. (1,3,2)
- B. (4,2,6)
- C. (4,3,1)
- D. All of the above
- E. None of the above



Clicker Question: Grouping

```
SELECT a1.x, a2.y, COUNT(*)  
FROM Arc a1, Arc a2  
WHERE a1.y = a2.x  
GROUP BY a1.x, a2.y
```

Arc a1

| x | y |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 2 | 3 |
| 3 | 4 |
| 3 | 4 |
| 4 | 1 |
| 4 | 1 |
| 4 | 1 |
| 4 | 2 |

Arc a2

| x | y |
|---|---|
| 1 | 2 |
| 1 | 2 |
| 2 | 3 |
| 3 | 4 |
| 3 | 4 |
| 4 | 1 |
| 4 | 1 |
| 4 | 1 |
| 4 | 2 |

Which of the following is in the result?

- A. (1,3,2)
- B. (4,2,6)
- C. (4,3,1)
- D. All of the above
- E. None of the above



Groupies of your very own

Find the average age for each standing (e.g., Freshman).

```
SELECT standing, avg(age)
FROM Student
GROUP BY standing
```



Groupies of your very own

Find the deptID and # of faculty members for each department having a department id > 20

Option 1:

```
SELECT count(*), deptid  
FROM faculty  
WHERE deptid > 20  
GROUP BY deptid
```

Option 2:

```
SELECT count(*), deptid  
FROM faculty  
GROUP BY deptid  
HAVING deptid > 20
```

Which one works?

A: Just 1

B: Just 2

C: Both

D: Neither